

Khwaja Mohammed Junaid Shaik

khwajamohammedjunaidshaik@gmail.com | +91 7337248619 | linkedin.com/in/junaid7623 | github.com/gyrfalcon55

SUMMARY

B.Tech (CSE – Data Science) student experienced in designing and deploying Generative AI and Agentic AI solutions using Python. Built production-grade LLM pipelines, RAG systems, explainable AI, and real-time inference APIs with prompt engineering and cloud deployment on Azure and AWS. Skilled in model training, feature engineering, experiment tracking, and backend integration via FastAPI and Flask, with a strong interest in scalable, cloud-based AI systems for business impact.

EDUCATION

B.Tech Computer Science & Engineering (Data Science)	2022 - 2026
Raghu Engineering College	CGPA: 8.1
Intermediate (MPC)	2020 - 2022
Narayana Junior College	73%
10th	2019 - 2020
Madhurawada English Medium School	92%

SKILLS

Languages: Python, SQL(PostgreSQL)

Machine Learning / Deep Learning: TensorFlow, Keras, Scikit-learn, NLTK

GenAI: LangChain, LangGraph, LangSmith, RAG, Prompt Engineering, OpenAI API, Groq API, Hugging Face

Backend Development: FastAPI, Flask

Data handling & Visualization: Pandas, NumPy, Matplotlib

Developer Tools: Git/GitHub, GitHub Actions, Docker, Mlflow

Cloud: AWS (EC2, S3, Lambda, IAM), Azure

PROJECT EXPERIENCE

EfficientNet-Based Skin Lesion Diagnosis with Explainable AI & Report Generation

<https://github.com/gyrfalcon55/Skin-Lesion-Diagnosis-and-Report-Generation>

- Developed a deep learning-based skin lesion classification system using EfficientNet and the HAM10000 dataset to classify seven types of skin diseases.
- Applied image preprocessing and lesion segmentation techniques to improve model performance, achieving **89% weighted F1-score**.
- Integrated Grad-CAM visualizations to highlight image regions influencing model predictions for better interpretability.
- Generated diagnostic summaries using the LangChain and OpenAI API based on model predictions.
- The application was deployed in AzureCloud, using Streamlit for real-time prediction and report generation.
- Tech Stack:** Python, TensorFlow, Keras, OpenCV, Flask, OpenAI API, LangChain

Legal Contract Processing & Clause Extraction Pipeline

https://github.com/gyrfalcon55/Document_Processing_with_LLMs

- Built an RAG-based legal document processing pipeline using LangChain and ChromaDB for contract understanding and clause extraction.
- Implemented a semantic search to retrieve relevant contract sections from **50+ diverse contract documents**.
- Used locally hosted LLMs through Ollama to generate structured responses for clauses such as termination, confidentiality, and liability.
- Designed prompt templates to generate responses based on retrieved contract content.
- Developed the pipeline to work across multiple contract formats and document structures.
- Tech Stack:** Python, LangChain, ChromaDB, Ollama(Llama 3.2, nomic-embed-text), Pydantic, RAG, Prompting

IMDB Sentiment Analysis - End-to-End NLP Pipeline with MLflow & FastAPI

https://github.com/gyrfalcon55/IMDB_Sentiment_Analysis

- Developed an end-to-end NLP pipeline for sentiment analysis on 50,000 IMDB movie reviews using Scikit-learn and NLTK.
- Implemented text preprocessing, feature extraction, and trained multiple machine learning models including Logistic Regression, Naive Bayes, and SVM.
- Performed hyperparameter tuning using RandomizedSearchCV and tracked experiments using MLflow and DagsHub.
- Built a FastAPI application for real-time sentiment prediction through a web interface.
- Integrated DVC for dataset versioning and reproducible workflow management.
- Tech Stack:** Python, Scikit-learn, NLTK, FastAPI, Docker, Git

CERTIFICATIONS

Oracle Cloud Infrastructure 2025 Certified AI Foundations Associate (link)